

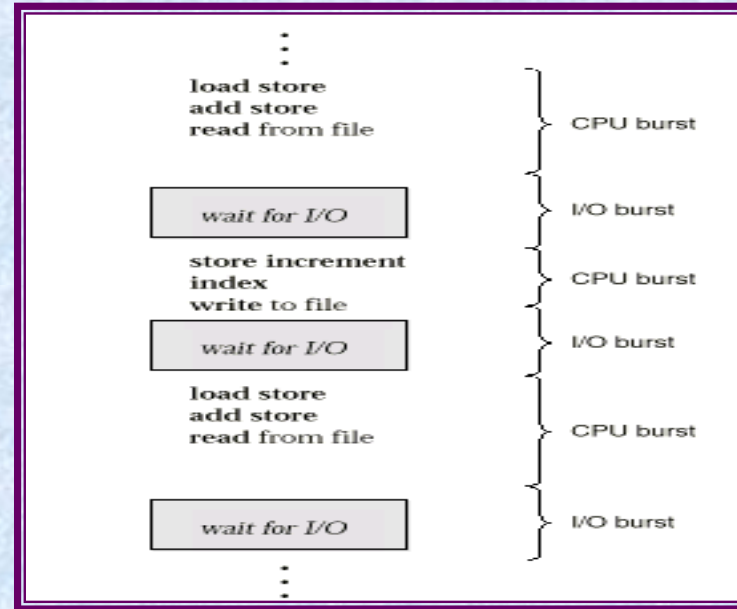
Chapter 6: CPU Scheduling

- Basic Concepts
- Scheduling Criteria
- Scheduling Algorithms
- Multiple-Processor Scheduling
- Real-Time Scheduling
- Algorithm Evaluation

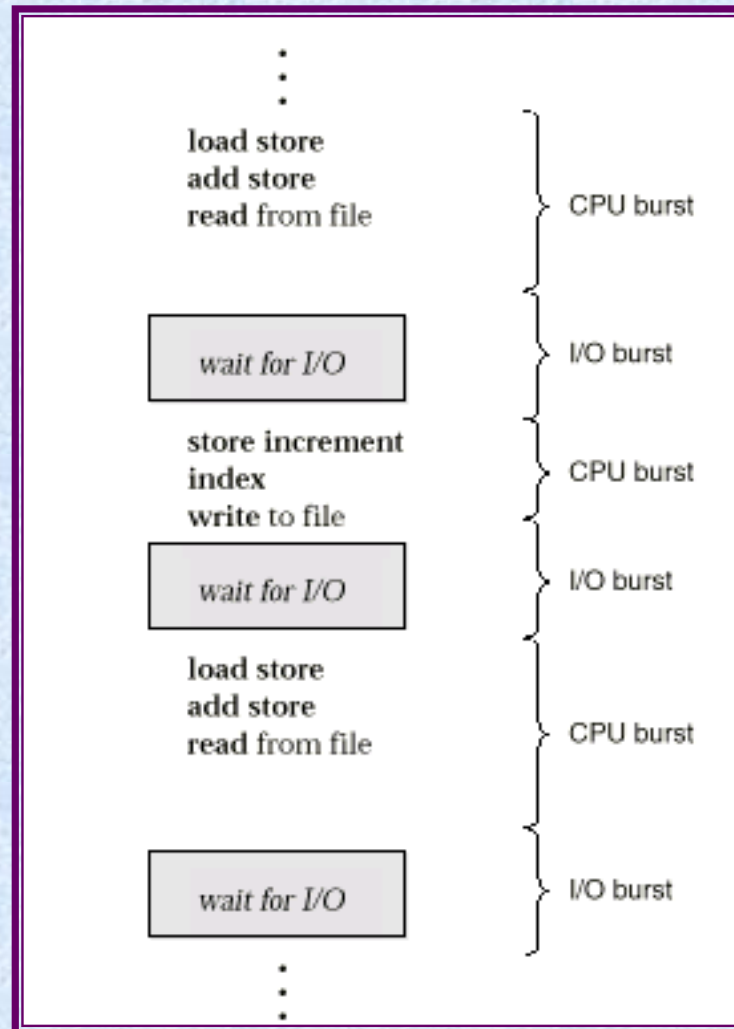
Basic Concepts

- Maximum CPU utilization obtained with multiprogramming
- CPU-I/O Burst Cycle – Process execution consists of a *cycle* of CPU execution and I/O wait.
- CPU burst distribution


Alternating Sequence of CPU And I/O Bursts



Alternating Sequence of CPU And I/O Bursts



Histogram of CPU-burst Times



CPU Scheduler

- Selects from among the processes in memory that are ready to execute, and allocates the CPU to one of them.
- CPU scheduling decisions may take place when a process:
 1. Switches from running to waiting state.
 2. Switches from running to ready state.
 3. Switches from waiting to ready.
 4. Terminates.
- Scheduling under 1 and 4 is *nonpreemptive*.
- All other scheduling is *preemptive*.

Dispatcher

- Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
 - ☞ switching context
 - ☞ switching to user mode
 - ☞ jumping to the proper location in the user program to restart that program
- *Dispatch latency* – time it takes for the dispatcher to stop one process and start another running.

Scheduling Criteria

- CPU utilization – keep the CPU as busy as possible
- Throughput – # of processes that complete their execution per time unit
- Turnaround time – amount of time to execute a particular process
- Waiting time – amount of time a process has been waiting in the ready queue
- Response time – amount of time it takes from when a request was submitted until the first response is produced, **not** output (for time-sharing environment)


Optimization Criteria

- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time

First-Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- Suppose that the processes arrive in the order: P_1, P_2, P_3
The Gantt Chart for the schedule is:




- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: $(0 + 24 + 27)/3 = 17$

FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order

$$P_2, P_3, P_1.$$

- The Gantt chart for the schedule is:



- Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- Average waiting time: $(6 + 0 + 3)/3 = 3$
- Much better than previous case.
- *Convoy effect* short process behind long process


Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time.
- Two schemes:
 - ☞ nonpreemptive – once CPU given to the process it cannot be preempted until completes its CPU burst.
 - ☞ preemptive – if a new process arrives with CPU burst length less than remaining time of current executing process, preempt. This scheme is known as the Shortest-Remaining-Time-First (SRTF).
- SJF is optimal – gives minimum average waiting time for a given set of processes.

Example of Non-Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

■ SJF (non-preemptive)




■ Average waiting time = $((0-0) + (8-2) + (7-4) + (12-5))/4$
 $= 0 + 6 + 3 + 7 = 16/4 = 4$

■ Average Turn time = $((0+7) + (6+4) + (3+1) + (7+4))/4$
» $7 + 10 + 4 + 11 = 32 / 4 = 8$

Example 1 of Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

■ SJF (preemptive)



■ Average waiting time = $\frac{((0-0)+(11-2)) + ((2-2)+(5-4)) + (4-4) + (7-5)}{4}$
 $= \frac{(0 + 9) + (0 + 1) + (0) + (2)}{4} = \frac{12}{4} = 3$

■ Average TAT = $\frac{((9+7) + (1+4) + (0+1) + (2+4))}{4} = \frac{16+5+1+6}{4} = \frac{28}{4} = 7$

Example 2 Preemptive SJF


Process	Arrival Time	Burst Time
---------	--------------	------------

P_1	0.0	3
-------	-----	---

P_2	2.0	4
-------	-----	---

P_3	4.0	1
-------	-----	---

P_4	5.0	2
-------	-----	---



- Average waiting time = $(0 + (3 - 2) + (7 - 4) + (4 - 4) + (5 - 5)) = 1 + 3 + 0 + 0 = 4 / 4 = 1$
- Avg TaT = $(0 + 3) + (4 + 4) + (0 + 1) + (0 + 2) = 3 + 8 + 1 + 2 = 14 / 4 = 3.5$

Determining Length of Next CPU Burst

- Can only estimate the length.
- Can be done by using the length of previous CPU bursts, using exponential averaging.

1. t_n = actual length of n^{th} CPU burst
2. τ_{n+1} = predicted value for the next CPU burst
3. $\alpha, 0 \leq \alpha \leq 1$
4. Define :


$$\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n.$$

Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority).
 - ☞ Preemptive
 - ☞ nonpreemptive
- SJF is a priority scheduling where priority is the predicted next CPU burst time.
- Problem \equiv Starvation – low priority processes may never execute.
- Solution \equiv Aging – as time progresses increase the priority of the process.

Non-Preemptive Priority


Process	Arrival Time	Burst Time	Priority
P_1	0.0	7	3
P_2	2.0	4	4
P_3	4.0	1	1
P_4	5.0	4	2



- Wt. Time $P_1 = 0 - 0 = 0$, $P_2 = 12 - 2 = 10$,
 $P_3 = 7 - 4 = 3$, $P_4 = 8 - 5 = 3$
- Average waiting time = $(0 + (12.0 - 2.0) + (7.0 - 4.0) + (8.0 - 5.0)) / 4 = 10 + 3 + 3 = 16 / 4 = 4$
- Turn Around Time $P_1 = 0 + 7 = 7$, $P_2 = 10 + 4 = 14$
 $P_3 = 3 + 1 = 4$, $P_4 = 3 + 4 = 7$
- Avg TaT = $(7 + 14 + 4 + 7) / 4 = 32 / 4 = 8$

Preemptive Priority Example

Process	Arrival Time	Burst Time	Priority
<i>P1</i>	0.0	7	3
<i>P2</i>	2.0	4	4
<i>P3</i>	4.0	1	1
<i>P4</i>	5.0	4	2



- Wt. Time $P1 = (0-0) + (9-4) = 5$, $P2 = 12 - 2 = 10$,
 $P3 = 4 - 4 = 0$, $P4 = 5 - 5 = 0$
Avg. wt. time = $5 + 10 + 0 + 0 = 15/4 = 3.75$
- TAT Time $P1 = 5 + 7 = 12$, $P2 = 10 + 4 = 14$,
 $P3 = 0 + 1 = 1$, $P4 = 0 + 4 = 4$
Avg TAT = $(12 + 14 + 1 + 4)/4 = 31/4 = 7.75$


Round Robin (RR)

- Each process gets a small unit of CPU time (*time quantum*), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once. No process waits more than $(n-1)q$ time units.
- Performance
 - ☞ q large \Rightarrow FIFO
 - ☞ q small $\Rightarrow q$ must be large with respect to context switch, otherwise overhead is too high.

Example of RR with Time Quantum = 20

<u>Process</u>	<u>Burst Time</u>
P_1	53
P_2	17
P_3	68
P_4	24


- The Gantt chart is:



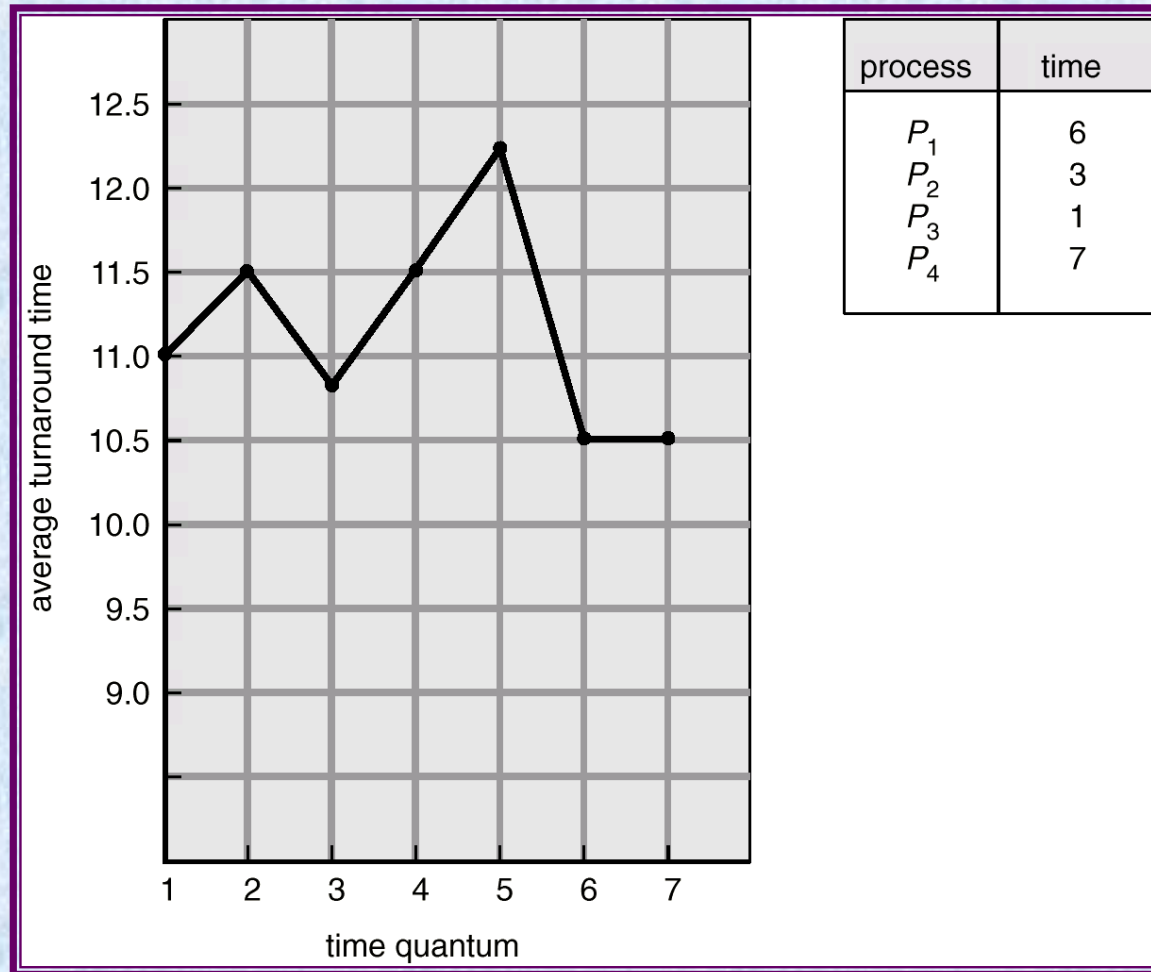
- Wt. Time $P_1 = (0-0) + (77-20) + (121-97) = 81$, $P_2 = 20-0 = 20$,
 $P_3 = (37-0) + (97-57) + (134-117) = 94$, $P_4 = (57-0) + (117-77) = 97$

- Avg. wt. time = $81+20+94+97 = 292/4 = 73$
- Typically, higher average turnaround than SJF, but better *response*.

Time Quantum and Context Switch Time



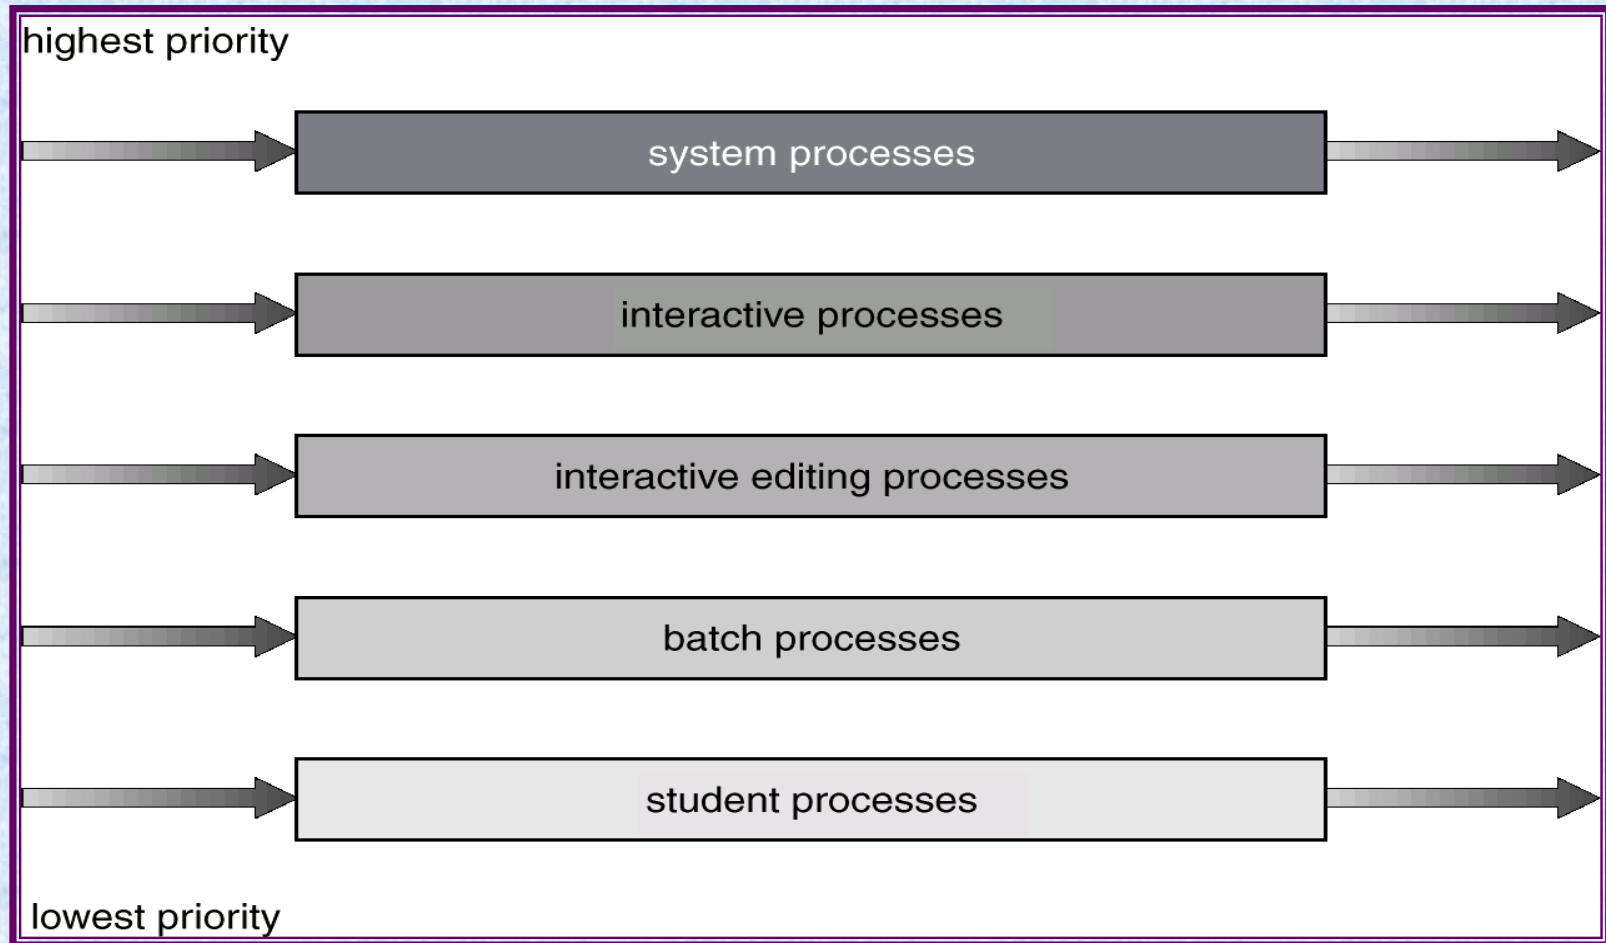
Turnaround Time Varies With The Time Quantum



Multilevel Queue

- Ready queue is partitioned into separate queues:
foreground (interactive)
background (batch)
- Each queue has its own scheduling algorithm,
foreground – RR
background – FCFS
- Scheduling must be done between the queues.
 - ☞ Fixed priority scheduling; (i.e., serve all from foreground then from background). Possibility of starvation.
 - ☞ Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., 80% to foreground in RR
 - ☞ 20% to background in FCFS

Multilevel Queue Scheduling



Multilevel Feedback Queue

- A process can move between the various queues; aging can be implemented this way.
- Multilevel-feedback-queue scheduler defined by the following parameters:
 - ☞ number of queues
 - ☞ scheduling algorithms for each queue
 - ☞ method used to determine when to upgrade a process
 - ☞ method used to determine when to demote a process
 - ☞ method used to determine which queue a process will enter when that process needs service

Example of Multilevel Feedback Queue


■ Three queues:

- ☞ Q_0 – time quantum 8 milliseconds
- ☞ Q_1 – time quantum 16 milliseconds
- ☞ Q_2 – FCFS

■ Scheduling

- ☞ A new job enters queue Q_0 which is served FCFS. When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue Q_1 .
- ☞ At Q_1 job is again served FCFS and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue Q_2 .

Multilevel Feedback Queues



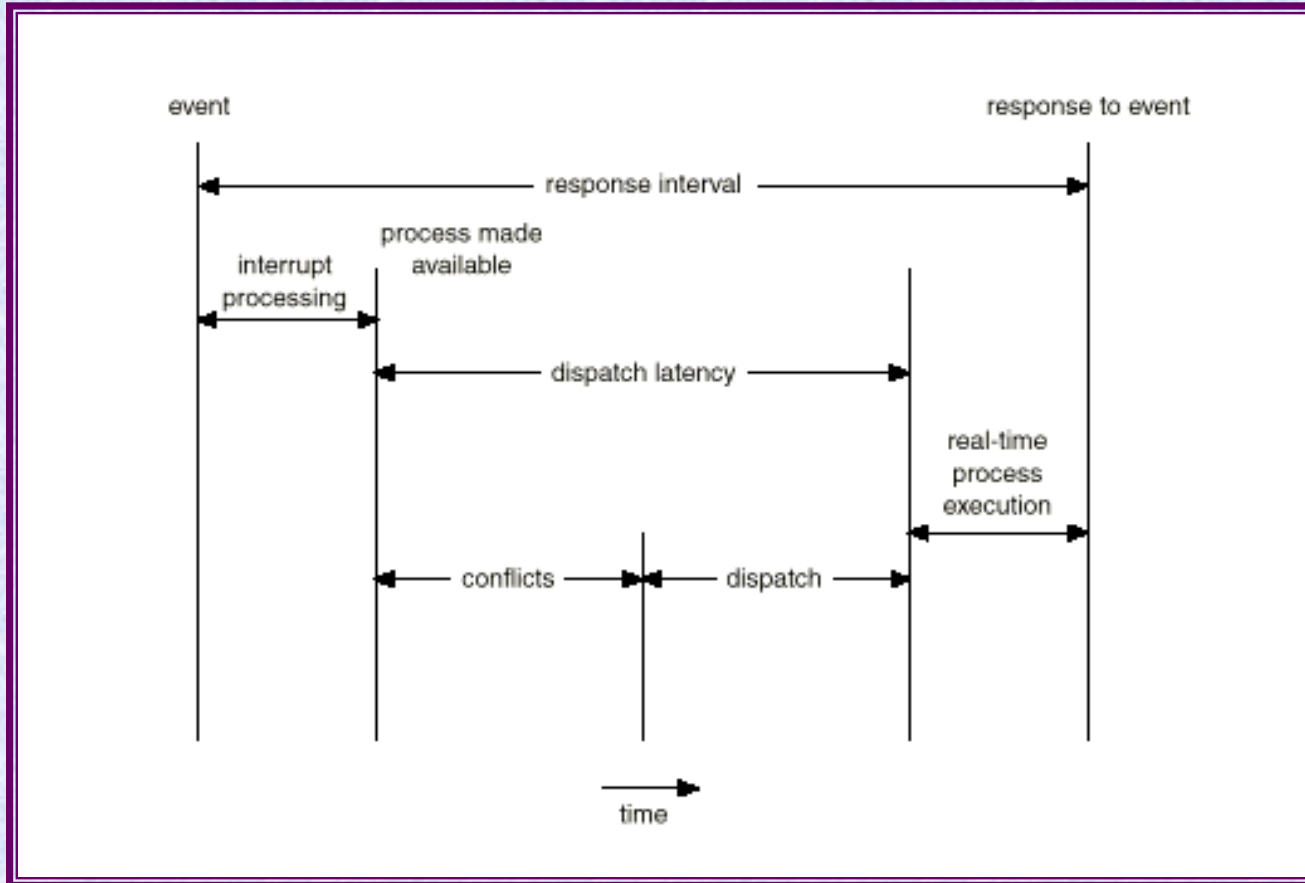
Multiple-Processor Scheduling

- CPU scheduling more complex when multiple CPUs are available.
- *Homogeneous processors* within a multiprocessor.
- *Load sharing*
- *Asymmetric multiprocessing* – only one processor accesses the system data structures, alleviating the need for data sharing.

Real-Time Scheduling

- *Hard real-time* systems – required to complete a critical task within a guaranteed amount of time.
- *Soft real-time* computing – requires that critical processes receive priority over less fortunate ones.

Dispatch Latency



Algorithm Evaluation

- Deterministic modeling – takes a particular predetermined workload and defines the performance of each algorithm for that workload.
- Queueing models
- Simulation
- Implementation

1.Deterministic modeling – takes a particular predetermined workload and defines the performance of each algorithm for that workload

§ .2.Queueing models - no static set of processes; determine the distribution of CPU & I/O bursts, mathematical formula describing the probability of a CPU burst; arrival time distribution.

- Computer system is described as network of servers.
- Each server has a queue of waiting processes.
- The CPU is a server with ready queue; I/O with device queue.
- Knowing the arrival rates and service rates; can compute CPU utilization, average queue length, average wait time, etc called as Queueing – network analysis.

Little's formula – - $n = \lambda \times w$ where n is average queue length, w is average wt. time in queue and λ is average arrival rate (3/sec).

3. Simulation –

- more accurate;
- Programming a model of the computer system

Software data structure represents the major component.

■ 4. Implementation –

- Most accurate
- Not always feasible
Expensive